

# Universal LLM Wrapper Documentation

## Overview

This document outlines the architecture and implementation of a **Universal LLM Wrapper System**, designed to interact with multiple AI model providers (OpenAI, Anthropic, Groq, DeepSeek, Fireworks, Together, Ollama, etc.) through a **centralized, extensible, and modular interface**.

The core goal is to:

- **Abstract provider-specific logic** from the application
- **Standardize request and response handling**
- **Simplify switching between models**
- **Support retry, timeout, and tool-use control**
- **Provide a universal CLI to test models**
- Enable integration with LangChain

---

## Folder Structure

```
llm_wrapper/  
├── providers_config.py      # Central configuration for all providers  
├── llm_client_factory.py    # Provider dispatcher (returns correct client)  
├── llm_cli.py               # CLI tool for model testing  
├── llm_clients/  
│   ├── __init__.py  
│   ├── base.py              # BaseLLMClient (abstract class)  
│   ├── openai_client.py  
│   ├── anthropic_client.py  
│   ├── groq_client.py  
│   ├── deepseek_client.py  
│   └── ...
```

---

## Core Concepts

### Why Use a Wrapper?

- Model APIs change, wrappers hide the volatility.
- Easily add/remove providers (like plugins).
- Central config (no need to touch app code).
- CLI/debug tools and UIs become simpler.

- Retry/timeout/config is unified.
- Abstracts LangChain or non-LangChain usage.

## providers\_config.py

Central registry for all LLMs, their models, capabilities, timeouts, retry logic, etc.

### Structure:

```
PROVIDERS = {
    "openai": {
        "name": "OpenAI",
        "base_url": "https://api.openai.com/v1",
        "models": {
            "gpt-4o": {
                "capabilities": ["chat", "stream", "json_mode",
"function_calling"],
                "default": True
            },
            "gpt-4o-mini": {
                "capabilities": ["chat", "stream", "json_mode"]
            }
        },
        "default_model": "gpt-4o-mini",
        "api_key_env": "OPENAI_API_KEY",
        "timeout": 30,
        "max_retries": 3
    },
    ...
}
```

### Capabilities Supported:

Capability	Meaning
chat	Supports chat completion APIs
stream	Supports streaming output
json_mode	Supports structured JSON mode generation
function_calling	Native function call support (OpenAI etc.)
tool_use	Advanced tool use support (Claude, Gemini)
code	Optimized for code (DeepSeek Coder, Codestral, etc.)

## API Key Handling

Each provider entry contains an `api_key_env` key which specifies the environment variable used to inject the key dynamically. This avoids hardcoding secrets in the codebase.

---

## `llm_client_factory.py`

This is the **dispatcher** that returns the correct client based on provider name.

```
def get_llm_client(provider_name, model=None, config=None):
    if provider_name == "openai":
        return OpenAIClient(model, config)
    elif provider_name == "anthropic":
        return AnthropicClient(model, config)
    ...
    else:
        raise ValueError(f"Unknown provider: {provider_name}")
```

---

## `llm_clients/base.py`

Defines a **standard interface** for all providers.

```
class BaseLLMClient:
    def __init__(self, model, config):
        self.model = model
        self.config = config

    def chat(self, messages, temperature=0.7):
        raise NotImplementedError
```

---

## Retry and Timeout

Use `tenacity` for retry logic:

```
from tenacity import retry, stop_after_attempt, wait_fixed

@retry(stop=stop_after_attempt(3), wait=wait_fixed(2))
```

---

```
def chat(...):  
    ...
```

The retry count and timeout duration should come from `config['max_retries']` and `config['timeout']`.

## Unified Response Format

Every `.chat()` method must return this schema:

```
{  
    "provider": "openai",  
    "model": "gpt-4o",  
    "message": "string response here",  
    "usage": {  
        "prompt_tokens": 45,  
        "completion_tokens": 130,  
        "total_tokens": 175  
    },  
    "raw": {...} # Optional raw API response  
}
```

This makes it easy to log, display in UI, debug, or do evals across providers.

## LangChain Compatibility

You can optionally power each client using LangChain wrappers internally:

```
from langchain.chat_models import ChatOpenAI  
  
class OpenAIClient(BaseLLMClient):  
    def chat(self, messages, temperature=0.7):  
        llm = ChatOpenAI(...)  
        response = llm.invoke(messages)  
        return {  
            "provider": "openai",  
            "model": self.model,  
            "message": response.content,  
            "raw": response  
        }
```

If the model is not available in LangChain (like Ollama or DeepSeek), you can just call HTTP directly.

## CLI Tool: llm\_cli.py

A test utility to call any provider/model from the command line.

### Usage:

```
$ python llm_cli.py --provider openai --model gpt-4o --prompt "Tell me a joke"
```

### Code:

```
import argparse
from providers_config import PROVIDERS
from llm_client_factory import get_llm_client




def main():
    parser = argparse.ArgumentParser()
    parser.add_argument("--provider", required=True)
    parser.add_argument("--model")
    parser.add_argument("--prompt", required=True)
    args = parser.parse_args()

    config = PROVIDERS[args.provider]
    model = args.model or config["default_model"]

    client = get_llm_client(args.provider, model, config)
    response = client.chat([{"role": "user", "content": args.prompt}])
    print(response["message"])

if __name__ == "__main__":
    main()
```

## Future-Proofing Ideas

Feature	Status
Eval multiple models	 Planned
Auto fallback	 Planned
Streaming via WebSocket	 Optional

Feature	Status
LangGraph Integration	Optional
Model pricing estimation	Optional
Token cost calculator	Optional
Usage caching	Optional
Unified telemetry/logging	Optional

---

## Conclusion

This wrapper gives you a powerful and flexible base to:

- Standardize model usage
- Scale to new providers quickly
- Create internal tooling (CLI, evals, dashboards)
- Swap LangChain in or out
- Keep all model logic in one place

It's modular, robust, and dev-friendly. Use this as your LLM DevOps foundation ✓